

06- Interrupt Handling

Ahmad Abdullah, PhD

abdullah@uvic.ca

<https://web.uvic.ca/~abdullah/csc230>

Lectures: MR 10:00 – 11:20 am


Location: ECS 125

Outline

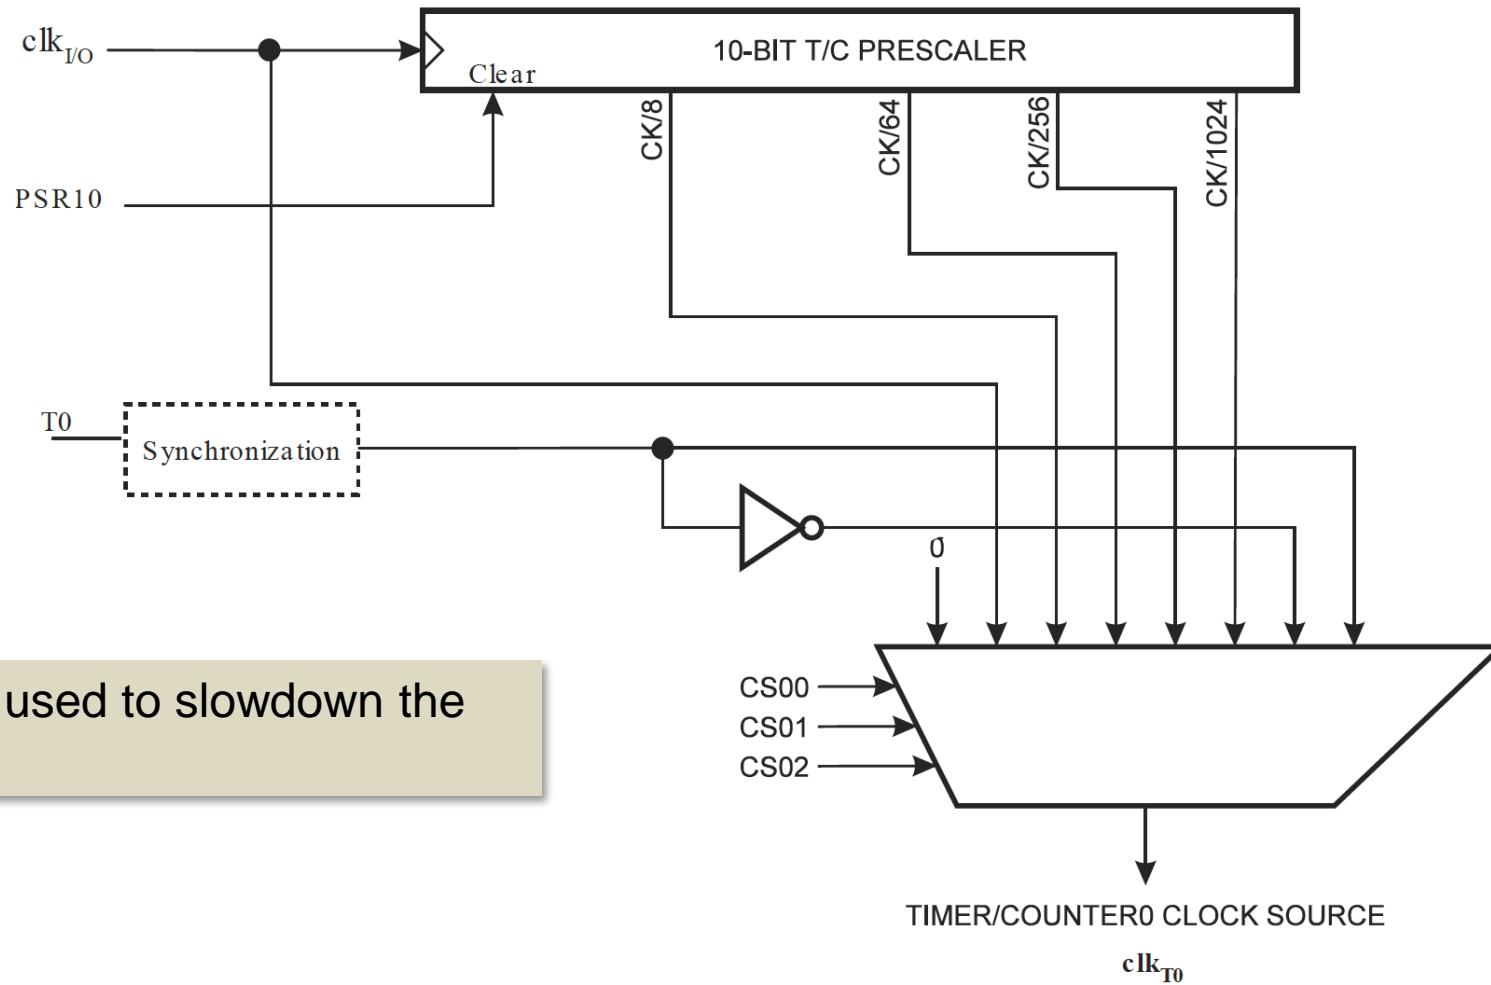
- Before interrupts: polling
- The interrupt concept
- Implementing interrupt handlers
- Nested interrupts
- More implementation issues

Timers

- A timer is a hardware device that keeps a track of an elapsed time between two events.
- A timer is simply a counter that counts the number of clock pulses and signals the CPU when it reaches a predefined count.
- The Atmel AT2560 processor supports several timers
 - Some of the timers are eight-bit timers (counters 0 & 2)
 - Most, however, are 16-bit timers (counters 1, 3, 4, and 5)
- In effect, these timers are bits of logic on the CPU that act independently of the ALU
 - Although they depend upon the CPU clock for their source of their oscillations...
 - ... they continue to tick independently of other code written by the programmer.




Timer diagrams



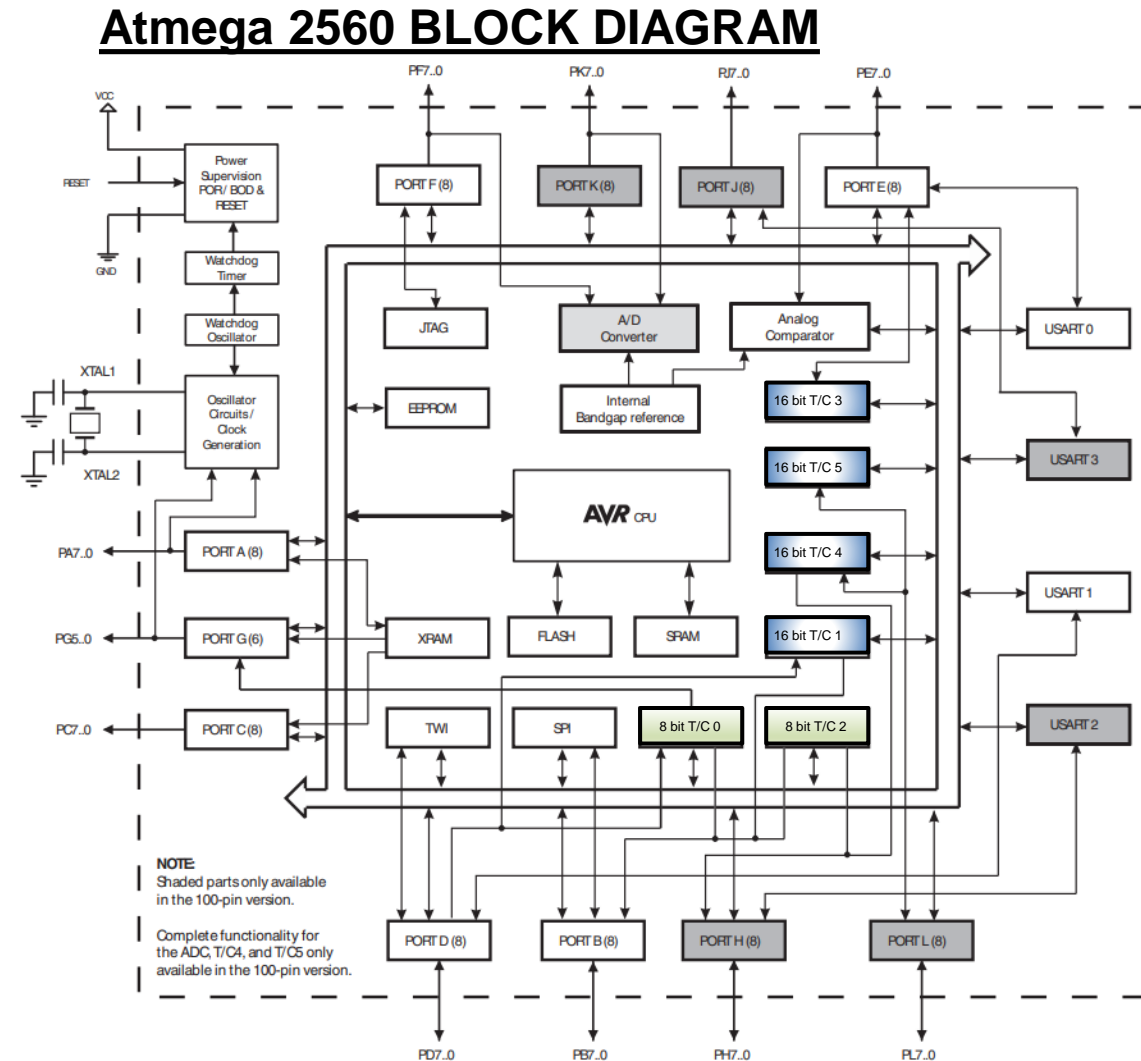
Prescalers are often used to slowdown the main clock.

Timer diagrams (Cont.)



Atmega 2560 Timers

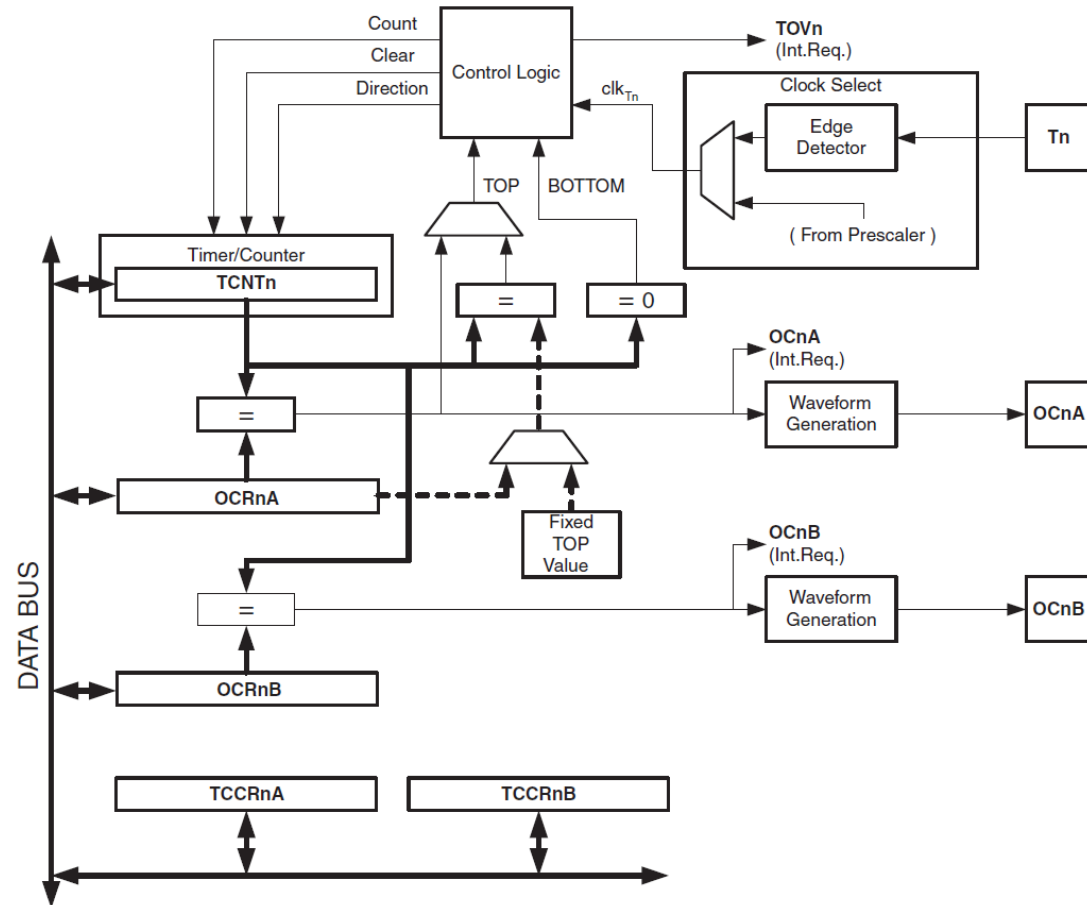
- The ATmega2560 has 6 timers/counters.
- Two 8-bit: T/C 0 and T/C 2.
- Four 16-bit: T/C 1, T/C 3, T/C 4 and T/C 5.
- Can be clocked by an internal or an external clock source.



8-bit Timer/Counter

Registers

- TCNT: Timer CouNTER register.
- OCR: Output Compare Registers.
- TCCR: Timer/Counter Control Register.
- OC: Output Compare Register.



TCCR2A – Timer/Counter Control Register A

7	6	5	4	3	2	1	0
COM2A1	COM2A0	COM2B1	COM2B0	–	–	WGM21	WGM20
R/W	R/W	R/W	R/W	R	R	R/W	R/W

Mode	WGM2	WGM1	WGM0	Timer/Counter Mode of Operation	TOP	Update of OCRx at	TOV Flag Set on
0	0	0	0	Normal	0xFF	Immediate	MAX
1	0	0	1	PWM, Phase Correct	0xFF	TOP	BOTTOM
2	0	1	0	CTC	OCRA	Immediate	MAX
3	0	1	1	Fast PWM	0xFF	BOTTOM	MAX
4	1	0	0	Reserved	–	–	–
5	1	0	1	PWM, Phase Correct	OCRA	TOP	BOTTOM
6	1	1	0	Reserved	–	–	–
7	1	1	1	Fast PWM	OCRA	BOTTOM	TOP

COM2A1	COM2A0	Description
0	0	Normal port operation, OC2A disconnected
0	1	Toggle OC2A on Compare Match
1	0	Clear OC2A on Compare Match
1	1	Set OC2A on Compare Match


TCCR2B – Timer/Counter Control Register B

7	6	5	4	3	2	1	0
FOC2A	FOC2B	–	–	WGM22	CS22	CS21	CS20
W	W	R	R	R/W	R/W	R/W	R/W

CS22	CS21	CS20	Description
0	0	0	No clock source (Timer/Counter stopped)
0	0	1	clk I/O/(No prescaling)
0	1	0	clk I/O/8 (From prescaler)
0	1	1	clk I/O/32 (From prescaler)
1	0	0	clk I/O/64 (From prescaler)
1	0	1	clk I/O/128 (From prescaler)
1	1	0	clk I/O/256 (From prescaler)
1	1	1	clk I/O/1024 (From prescaler)

TIMSK2 – Timer/Counter Interrupt Mask Register

- Bit 2 – OCIE2B: Timer/Counter Output Compare Match B Interrupt Enable
- Bit 1 – OCIE2A: Timer/Counter0 Output Compare Match A Interrupt Enable
- Bit 0 – TOIE2: Timer/Counter0 Overflow Interrupt Enable



Timers (Cont.)

- Our current method of using timers depends upon **polling**
- Polling is where:
 - Some device **external to the main CPU** (i.e., low-level hardware, timer, I/O interface) ...
 - ... has **some kind of state** that **must be checked by the CPU** for readiness ...
 - ... and for which **some action must be taken** .. (i.e., a value is ready to be input from a device; a device is ready to receive a value for output; the timer has reached its maximum value; etc.s) ...
 - ... and that state is checked -- and actions taken -- **via code contained within an infinite loop**.

Polling loop

```
# ...

codeToSetUpDevices()

# Polling loop
while True:
    if device[0].isReady():
        doSomethingForDevice(0)
    else if device[1].isReady():
        doSomethingForDevice(1)
    # etc etc
    else:
        someOtherNullStepPerhaps()

# We never expect to reach this point ...
```

Polling: pluses, minuses

- Polling is a relatively straightforward technique
 - We must still figure out how to write code to detect when devices are ready for activity ...
 - ... but that is not hard.
- As we add more devices or device events that must be managed by our code, we just add more **else if** clauses to our loop
 - In our timer example, we had two separate timers, each blinking at a different rate.
- **Two big downsides:**
 - Even if no device is ready for action, the CPU is still busy going around and around and around in the loop consuming CPU cycle
 - **Mixing polling code with non-polling code** can be difficult.

We can do better

- Working with low-level devices is a very important part of any computer system
- Writing code for such devices is always tricky
- So rather than polling...
 - i.e., our code constantly checking for a condition to be true ...
- ... we instead inform the computer that **events of interest to our program** should cause **one of our own functions** to be executed ...
- ... and that this function execution **must happen automatically**, regardless of what the rest of our code is doing at the time of the event!

What does it mean to be interrupted?




From: <https://www.youtube.com/watch?v=Mh4f9AYRCZY>

Interrupt

- **An event (signal) that can temporarily halt (“interrupt”) the control flow of the currently running program in order to execute a task on behalf of that event (signal).**
- These events can be triggered by:
 - external signals (e.g., external to the CPU such as I/O devices)
 - internal signals (e.g., internal to the CPU such as timers, CPU errors)
- The running program is temporarily suspended while the task is completed
 - This task usually is some function specially written for handling the event (i.e., interrupt handler)
 - The CPU takes care of transferring control to the interrupt handler and back to the interrupted running program.

Interrupt (Cont.)



A:

Device P generates an event, it is handled after a short delay

B:

Device Q generates an event, it is handled after a short delay

C:

Device P generates an event, while device Q generates an event when P's handle is executing.

Notice that any any one point in time, only a thick black line or thick green line is active.

Implicit / explicit

- The main code whose control-flow is suspended **does not know** an interrupt has occurred
 - That is, the main code does not call the interrupt-handler function
 - Transfer of control to handler is managed by the CPU
- However, the interrupt handler function **does explicitly cause control to return** to interrupted code
 - Using **reti** instead of ret

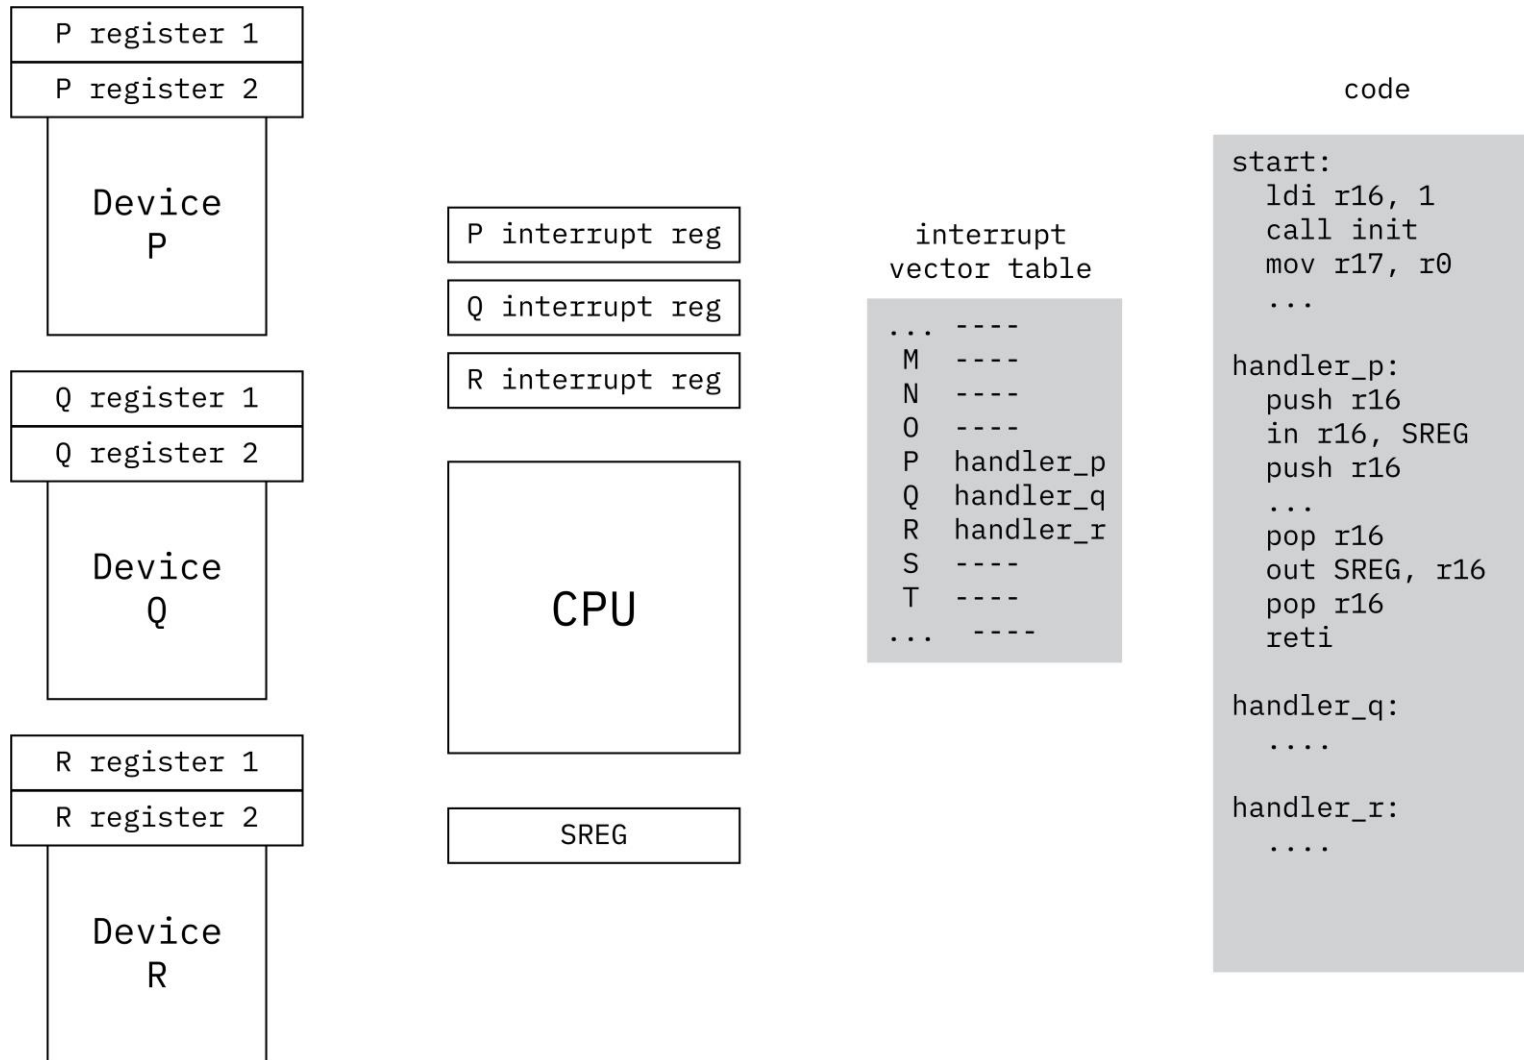
Interrupts can feel complicated

- Understanding how to program with interrupts can seem a bit tricky at first
- Some actions occur automatically, others require explicit actions on our part, and ultimately **many things will be happening simultaneously in the hardware** at runtime.
- Regardless:
 - Everything ultimately depends upon something we ourselves indicate in our code!
 - As long as we are clear about where we write code to configure interrupt handlers...
 - **... and as long as we very carefully write those handlers ...**
 - ... then coding can proceed in a straightforward fashion.

Overall idea

1. **With interrupts disabled on the CPU**, code that **configures the interrupt handler mappings** is itself executed
2. With this configuration complete, **enable interrupts**
3. Start to execute our main code
4. If an event occurs that is associated with an enabled interrupt...
 - a) ... the CPU pauses where we are in our code ...
 - b) ... transfers control to the correct handler for that event ...
 - c) ... and that handler disables interrupts ...
 - d) ... after which the handler does its work (i.e., code within the interrupt handler) ...
 - e) ... and when the handler is finished, the second-to-last thing it does is re-enable interrupts ...
 - f) ... and the last thing it does is causes the CPU to return control to where we originally paused.

Overall idea (Cont.)



Important

- Interrupts are not arbitrary programmatic things
 - They are not the same as threads in multithreaded programming
 - They are not used to implement multicore algorithms
- The CPU architecture will usually define the set of interrupts it can recognize
 - These interrupts are usually coupled to hardware devices the architects expect to be used
 - Some are to internal resources (such as timers)
- Programming for interrupts therefore means knowing something about what interrupt set the CPU supports
- ATmega2560: 57 interrupts (including the RESET interrupt)

Programming for interrupts

- **Interrupt vector entry:** entry in a special table that associates a device's interrupt with its interrupt-handler function
- **Interrupt handler:** A function that must be correctly implemented for the device
 - Requires knowledge of the device (obtained from a datasheet)
- Enabling some **specific interrupt flag**
 - This is for the actual device to permit it to raise interrupts
- Once all devices / interrupts are properly established...
 - ... only then enable global interrupts.
 - That is, until then these devices can be generating events and attempting to raise interrupts, but the CPU ignores them.
 - Only when **global interrupts are enabled** will the CPU transfer control to and from our handlers.

Interrupt Vectors

- The ATmega2560 **uses the first portion of program memory as the interrupt-vector table**
- Because we have not used interrupts so far in this course, **we haven't worried about this table...**
- **... but now we must!**
- To correctly make an entry we need to know:
 - The name of the interrupt
 - The interrupt-vector address at which the CPU expects the handler label
 - The handler label itself

Interrupt Vectors (Cont.)

14. Interrupts

This section describes the specifics of the interrupt handling as performed in ATmega640/1280/1281/2560/2561. For a general explanation of the AVR interrupt handling, refer to [“Reset and Interrupt Handling” on page 17](#).

14.1 Interrupt Vectors in ATmega640/1280/1281/2560/2561

Table 14-1. Reset and Interrupt Vectors

Vector No.	Program Address ⁽²⁾	Source	Interrupt Definition
1	\$0000 ⁽¹⁾	RESET	External Pin, Power-on Reset, Brown-out Reset, Watchdog Reset, and JTAG AVR Reset
2	\$0002	INT0	External Interrupt Request 0
3	\$0004	INT1	External Interrupt Request 1
4	\$0006	INT2	External Interrupt Request 2
5	\$0008	INT3	External Interrupt Request 3
6	\$000A	INT4	External Interrupt Request 4
7	\$000C	INT5	External Interrupt Request 5
8	\$000E	INT6	External Interrupt Request 6
9	\$0010	INT7	External Interrupt Request 7
10	\$0012	PCINT0	Pin Change Interrupt Request 0
11	\$0014	PCINT1	Pin Change Interrupt Request 1
12	\$0016 ⁽³⁾	PCINT2	Pin Change Interrupt Request 2
13	\$0018	WDT	Watchdog Time-out Interrupt

Interrupt Vectors (Cont.)

Table 14-1. Reset and Interrupt Vectors

Vector No.	Program Address ⁽²⁾	Source	Interrupt Definition
1	\$0000 ⁽¹⁾	RESET	External Pin, Power-on Reset, Brown-out Reset, Watchdog Reset, and JTAG AVR Reset
2	\$0002	INT0	External Interrupt Request 0
3	\$0004	INT1	External Interrupt Request 1
4	\$0006	INT2	External Interrupt Request 2

15	\$001C	TIMER2 COMPB	Timer/Counter2 Compare Match B
16	\$001E	TIMER2 OVF	Timer/Counter2 Overflow
17	\$0020	TIMER1 CAPT	Timer/Counter1 Capture Event
18	\$0022	TIMER1 COMPA	Timer/Counter1 Compare Match A
19	\$0024	TIMER1 COMPB	Timer/Counter1 Compare Match B
20	\$0026	TIMER1 COMPC	Timer/Counter1 Compare Match C
21	\$0028	TIMER1 OVF	Timer/Counter1 Overflow

```
.cseg
.org 0
    jmp start

.org 0x22
    jmp timer1_handler

...

start:
    ....

timer1_handler:
    ....
```

Interrupt Vectors (Cont.)

- We must be very careful here:
 - If we have turned on a device's interrupts...
 - ... but have not properly placed an entry in the vector table ...
 - ... then the ATmega2560 will use whatever values it finds at the table location (possibly random!) as the interrupt handler.
- Put differently:
 - The assembler **does not warn us if we have been sloppy with our programming.**
 - We must manually match interrupts used in our solution to handler labels written in the vector table.

Interrupt Handler

- Handlers are meant to be the **smallest amount of code necessary** to deal with the interrupt
 - For I/O devices, this would be reading bytes from, or writing bytes to, the registers of the device.
 - For timers, this would involve setting status variables to important values
- Code within an interrupt-handler function very rarely calls other functions
 - The time spent handling one interrupt is unavailable for handling those from other devices...
 - ... and given the time sensitivity of device behavior, missing another device's interrupt could be disastrous.
- Finally: handlers must ensure the interrupted code can safely resume operation after handler is finished

Example: reading button value

```
.....  
timer1_handler:  
    push r16  
    in r16, SREG  
    push r16  
  
    push XH  
    push XL  
    push r17  
    mov XH, high(BUTTON_COUNT)  
    mov XL, low(BUTTON_COUNT)  
    ld r17, X  
    inc r17  
    st X, r17  
    pop r17  
    pop XL  
    POP XH  
  
    pop r16  
    out SREG, r16  
    pop r16  
  
    reti
```

Save r16 for later restoration, and also save the current value of the status register.

Actual task meant for the handler.

Restore the status-register value, and also restore r16

Return from the interrupt.

Example: LED ON/OFF

```
.cseg
.def leds = r16 ;current LED state
.def temp = r18 ;a temporary register

.org 0; RESET
    jmp start

.org 0x1A; Timer 2 COMPA
    jmp LED_ON
.org 0x1C; Timer 2 COMPB
    jmp LED_OFF

.cseg
start:
    ldi temp, 0xFF ; set PORTB as output
    out DDRB, temp
    clr leds; To turn OFF all LEDs
    out PORTB, leds

    ; Clear general interrupts
    cli

    ; Clear timer 1 configurations
    clr temp
    sts TCCR2A, temp
    sts TCCR2B, temp
```

```
; Set to CTC (mode 2)
lds temp, TCCR2A
ori temp, (1<<WGM21)
sts TCCR2A, temp

; Set prescaler to /1024
lds temp, TCCR2B
ori temp, (1<<CS22) | (1<<CS21) |
                                     (1<<CS20)
sts TCCR2B, temp

; Set values for OCCR2A & OCR2B
ldi temp, 0xFF
sts OCR2A, temp

ldi temp, 0x7F
sts OCR2B, temp

; Enable interrupt A & B for timer 2
lds temp, TIMSK2
ldi temp, (1 << OCIE2A) | (1 << OCIE2B)
sts TIMSK2, temp

; Reset timer 2 (optional)
clr temp
sts TCNT2, temp
sei ; Enable general interrupts
```

```
end:
    rjmp end

LED_ON:
    push leds
    push temp
    in temp, SREG
    push temp
    ldi leds, 0xFF; To turn ON all LEDs
    out PORTB, leds
    pop temp
    out SREG, temp
    pop temp
    pop leds
    reti

LED_OFF:
    push leds
    push temp
    in temp, SREG
    push temp
    clr leds; To turn OFF all LEDs
    out PORTB, leds
    pop temp
    out SREG, temp
    pop temp
    pop leds
    reti
```

Enabling interrupts

- There are two levels here:
 - Enabling the interrupts **for the device itself**
 - Enabling the CPU **to respond to interrupts in general.**
- Sometimes the former is the same thing as “turning on the device”
 - We see this with the 16-bit timers (i.e., timer-interrupt mask register, such as TIMSK1)
 - Setting a bit on this register has the same effect as enabling interrupts to occur for that timer
- The former is the main switch and is a single operation: **sei** (set global interrupt flag)
- Therefore, usual format of code:
 - Set up all handlers, enable interrupts on all devices
 - Once this is done, then (and only then!) perform **sei**

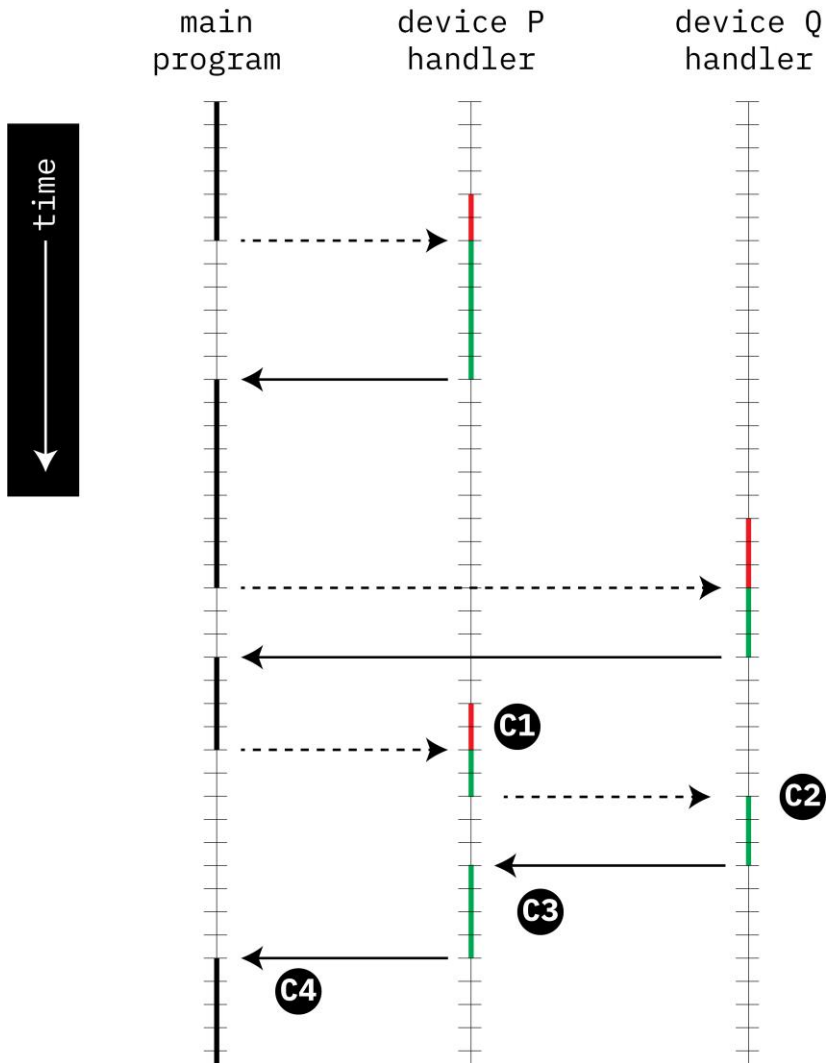
NMI

- In our handler-vector example, we saw one at location 0x00
- This vector is for the **RESET** signal
 - This is a special interrupt...
 - ... **which can never be ignored!**
 - It is also never meant something from which we return
 - Sometimes referred to as a “non-maskable interrupt” or “**NMI**”
- In our case, the vector-table entry just referred to the first location of our main program (the label start that we would have written)

Multi-level interrupts

- Our examples have also assumed that only one interrupt can be handled at any one time
 - In fact, when the CPU dispatches to an interrupt handler...
 - ... it clears the global interrupt flag in **SREG** ...
 - ... and only restores it after an **RETI**
- The ATmega2560 therefore supports **single-level interrupts**
 - All interrupts have the same priority or importance
- Other computer architectures support **multi-level interrupts**
 - Here a device and its interrupt has a priority...
 - ... such that interrupts from higher-priority devices are permitted to interrupt the handlers for lower-priority devices

Multi-level interrupts (Cont.)



If device Q had a higher priority than device P...

C:

Device P generates an event, but during its handling there is an event from device Q -- and so device Q (with higher priority) has its interrupt serviced.

Once that is done, the handlers for P is resumed.

Multi-level interrupts (Cont.)

- Hardware support is normally needed for multi-level interrupts
- Motivation for such an architecture:
 - Some devices generate many more events than others (i.e., SSD interface card vs. keyboard interface)
 - Some devices have data or state that is very volatile and must be quickly read (or written) within very closer tolerances.

Critical section code

- Sometimes we may write code for which interrupts might cause a problem
- Example:
 - Timer interrupts are often used by operating systems to force switching between processes.
 - However, code in the OS may want to avoid such a switch when executing some sensitive code within the OS kernel itself
 - Therefore, the OS will disable interrupts before that sensitive code, and re-enable them afterwards
- On ATmega2560: **cli** disables interrupts, **sei** enables interrupts.

Summary

- Interrupts permit us:
 - Write code that responds to devices running in parallel with the CPU ...
 - ... in such a way that servicing those devices does not require a polling loop.
- Concepts / vocabulary:
 - interrupt / interrupt event
 - interruption of main-program control flow (and need to push and pop the status registers)
 - handlers
 - interrupt vectors
 - setting up interrupts
 - enabling / disabling interrupts
 - nested interrupts



Any Questions?