

07-Performance Issues

Ahmad Abdullah, PhD

abdullah@uvic.ca

<https://web.uvic.ca/~abdullah/csc230>

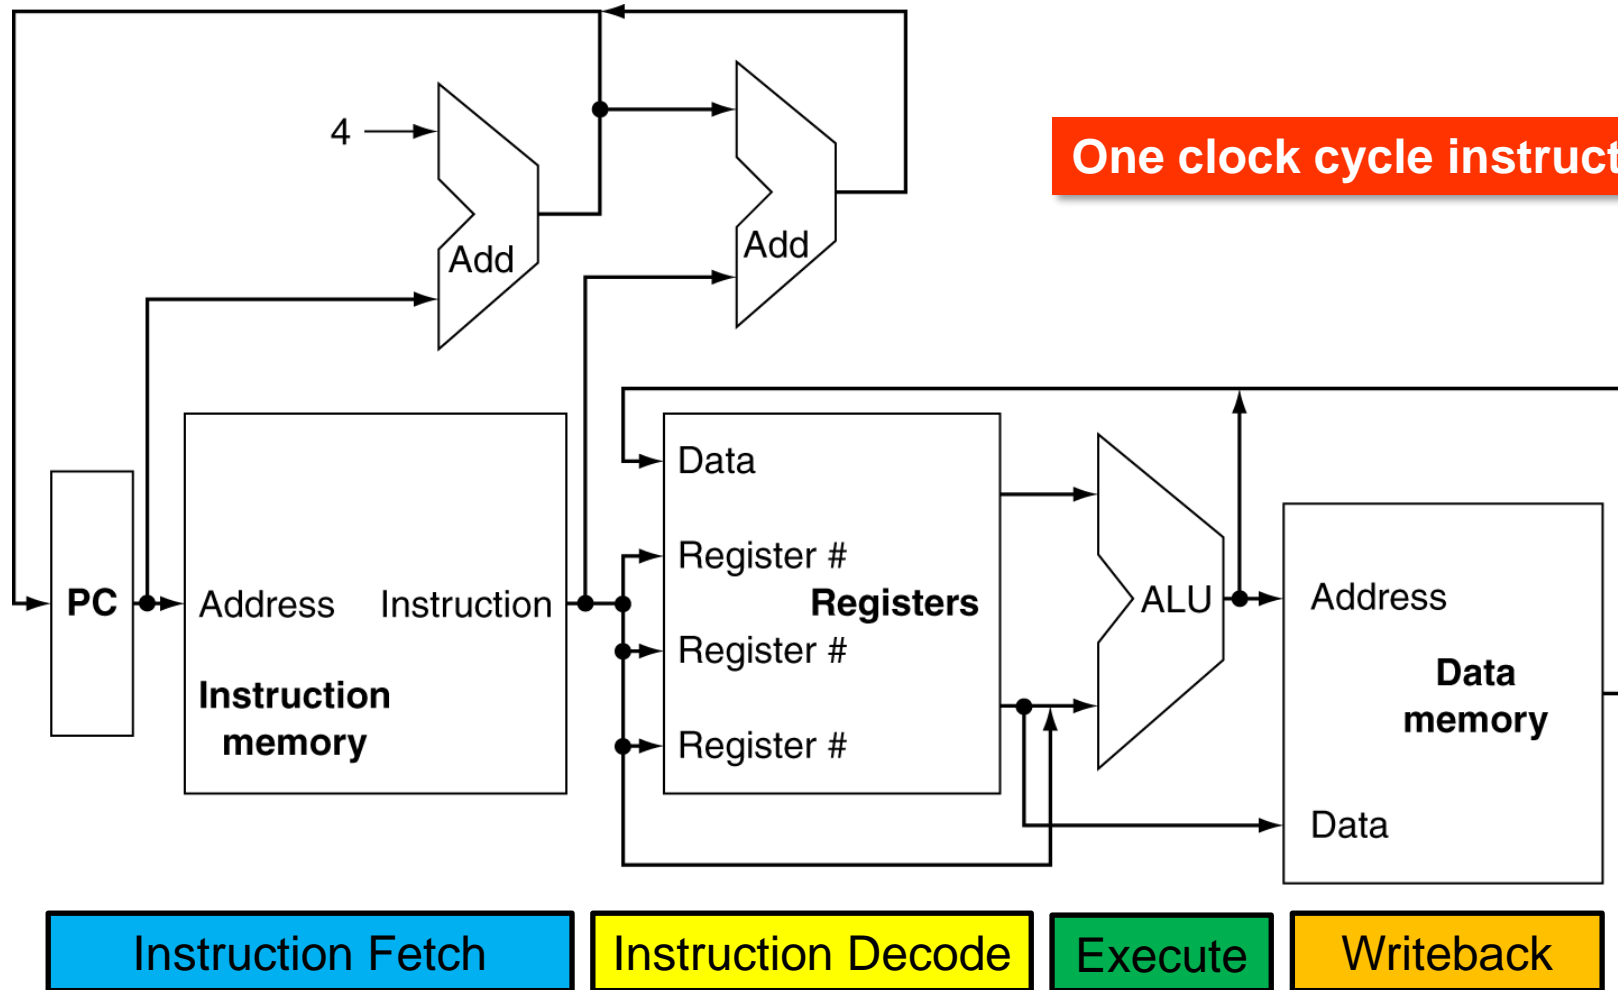
Lectures: MR 10:00 – 11:20 am

Location: ECS 125


Outline

- Pipelining
- Multicore
- Multithreading
- Amdahl's Law
- Measures of performance

Instruction Execution Cycle




Program Execution



Pipelining

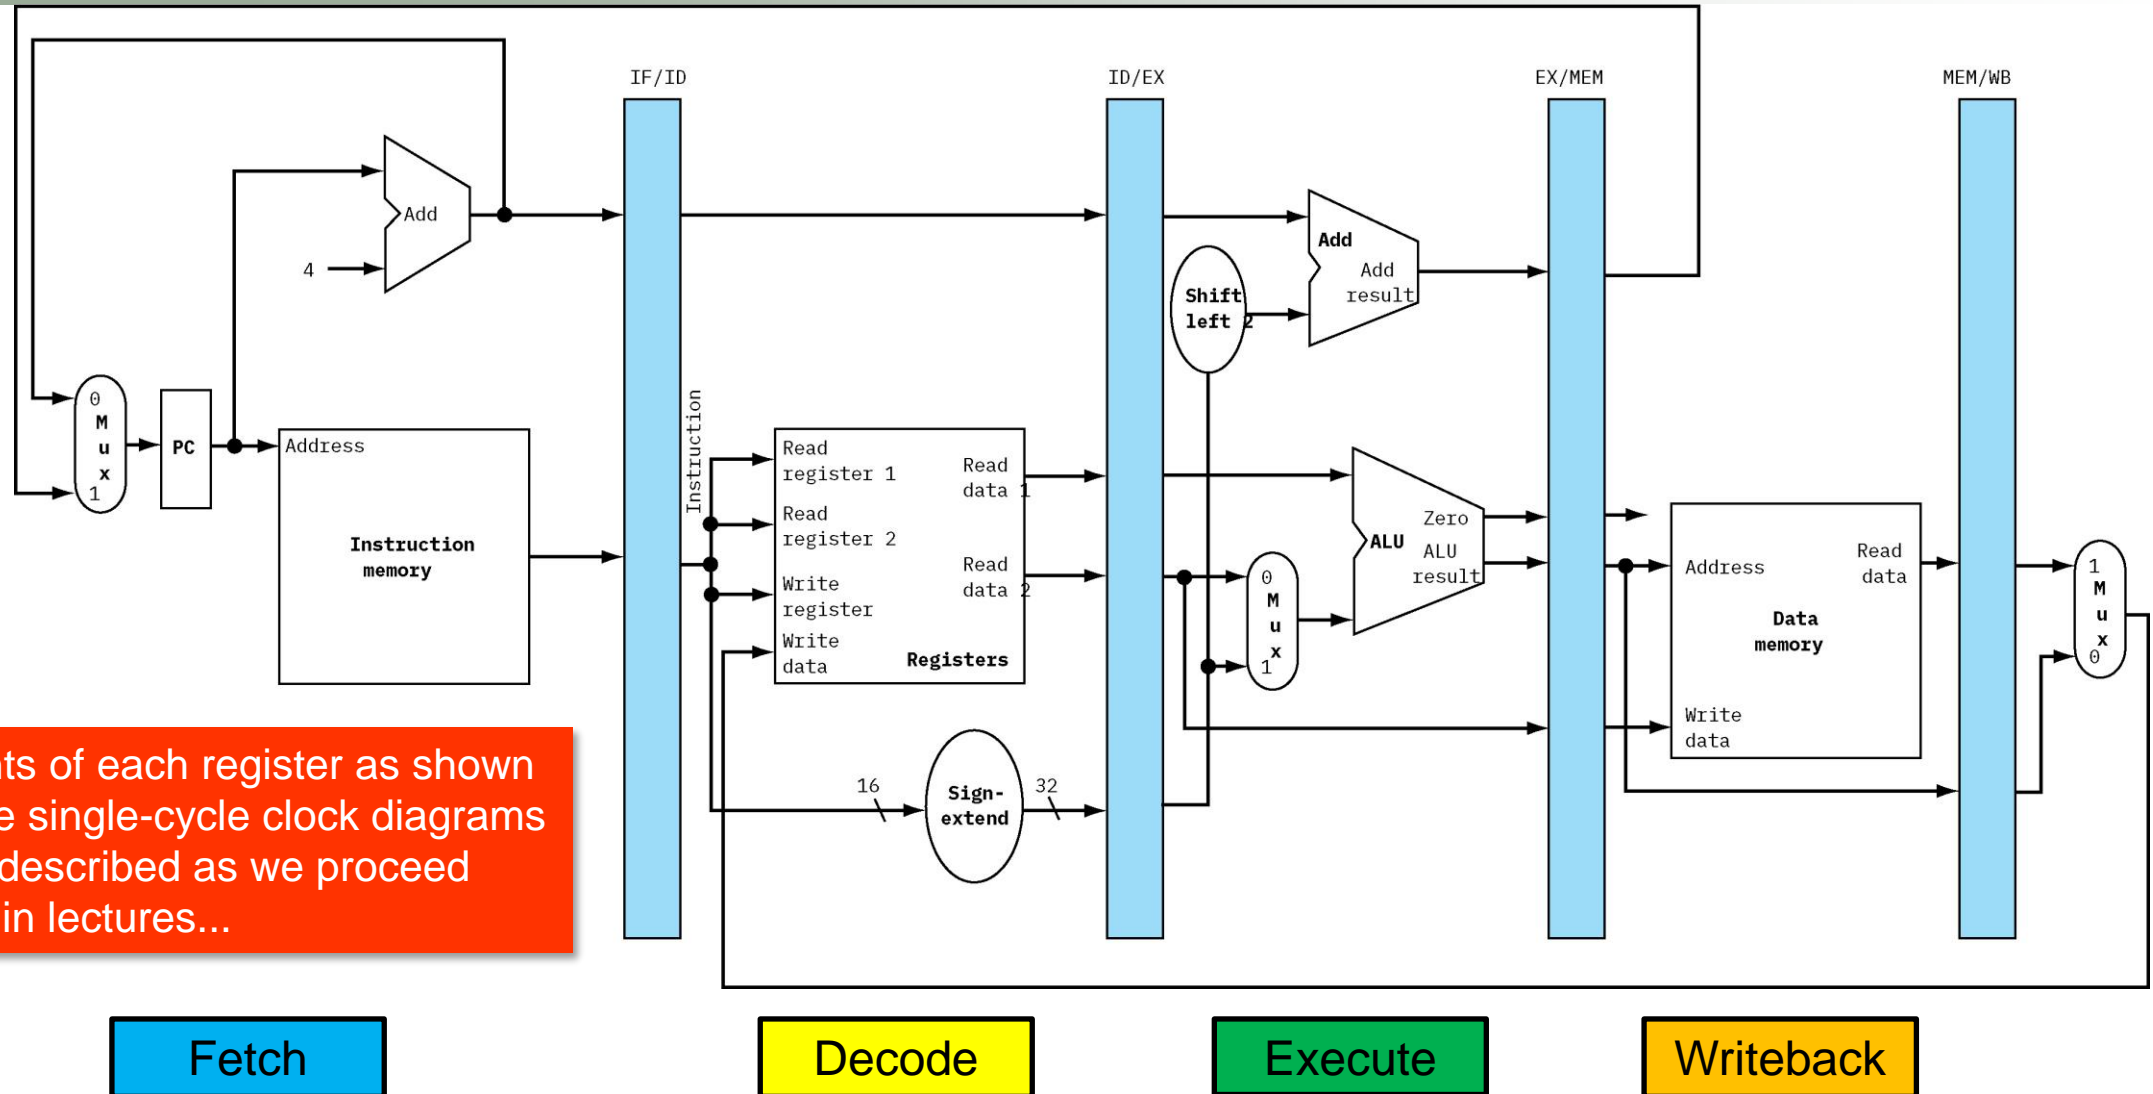
- Intuition:
 - like an assembly line
- Shrink cycle time (i.e., increase clock frequency)
- Now each phase fits within a clock cycle
- Different instructions at different stages can be worked on simultaneously (“instruction-level parallelism”)



During cycle 8:

- instruction 5 is in its final phase
- instruction 6 is in its execute phase
- instruction 7 is in its decode phase
- instruction 8 has just been fetched

Pipeline Stages




Pipeline Hazards

- Control hazard:
 - branching issue
- Structural hazard:
 - more than one instruction in the pipeline needs the CPU resource in the same cycle
- Data hazard:
 - instruction in earlier stage requires data that will be the result of a instruction in later stage (i.e., results only available after writeback phase)

Branching Issue

- If a branch is taken, then work for instructions still in the pipeline following the branch instruction is discarded
- Instruction 3 was a branch...
- ... and the branch is taken (i.e., to instruction 9)



Branching Prediction

- **Idea:** use actual program behavior to predict future behavior
- Processor keeps track of branch history
 - If CPU reaches branch instruction...
 - ... it looks up the instruction address in a table
 - ... then fetches the state of the branch
- If the state is “branch normally taken”
 - then target address fetched

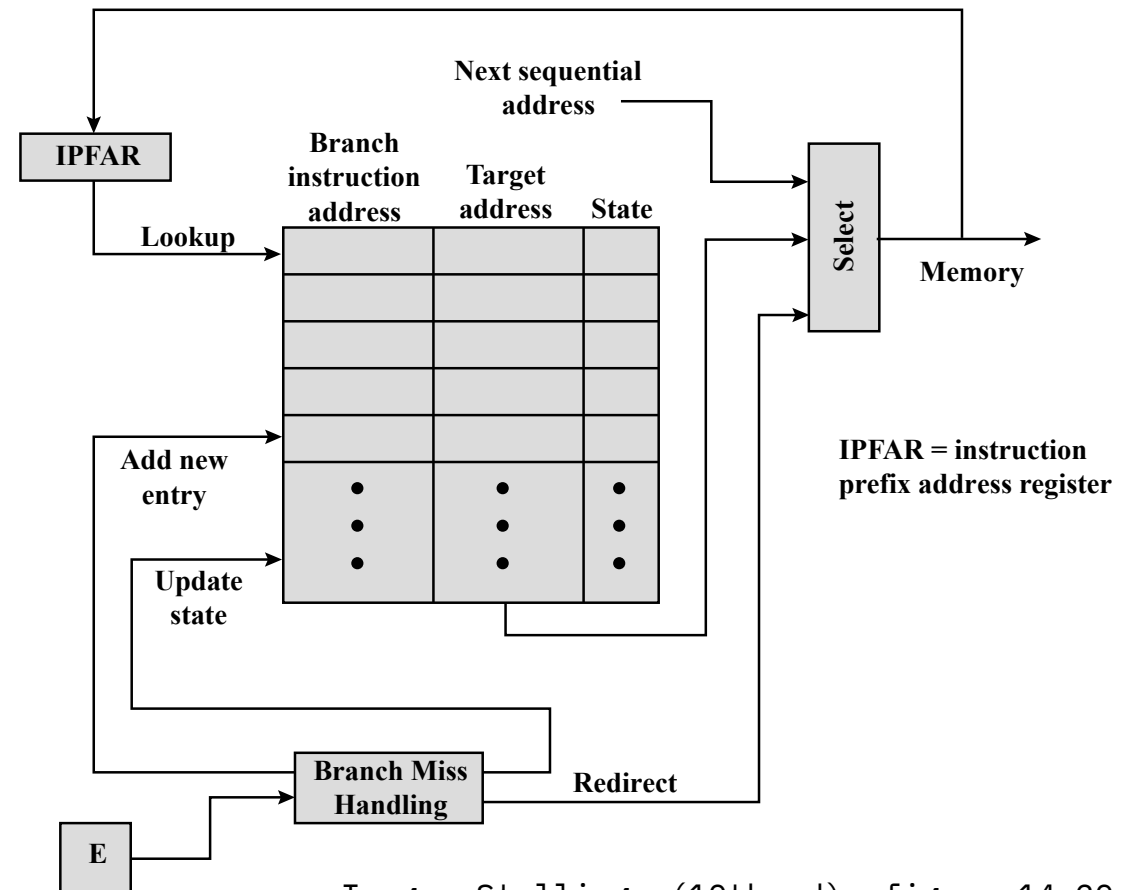


Image: Stallings (10th ed), figure 14.20

Branching Prediction (Cont.)

- If the branch is correctly predicted:
 - ... then the instructions fetched from memory based on that prediction (i.e., taken or not taken)...
 - ... and fed into the pipeline ...
 - ... will not result in wasted work.
- Of course, sometimes the prediction will be wrong
 - Much work has been done on increasing prediction accuracy
 - Example: using more and more history to generated branch prediction
 - (At some point, predictions must be wrong, otherwise we wouldn't need to have branches in the first place!)

Superscalar Execution

- This departs from basic pipelining in an important way.
- More than one instruction is fetched on the same clock cycle
 - In essence, multiple parallel pipelines




Image: Stallings (10th ed), figure 16.2

Data flow Analysis

- Processor determines possible equivalent orderings of instructions:
 - Ultimate goal: choose an order that executes the fastest
- **Data flow:** analyzing code to discover how the results of computations are used as inputs/operands for later computations

Original Code

```
a = b + c - d;  
if (e != 0) ...
```

Unscheduled

```
load  b, r1  
load  c, r2  
add    r1, r2, r3  
load  d, r4  
sub    r3, r4, r5  
store  r5, a  
load  e, r6  
cmp    r6, 0  
bne    ...
```

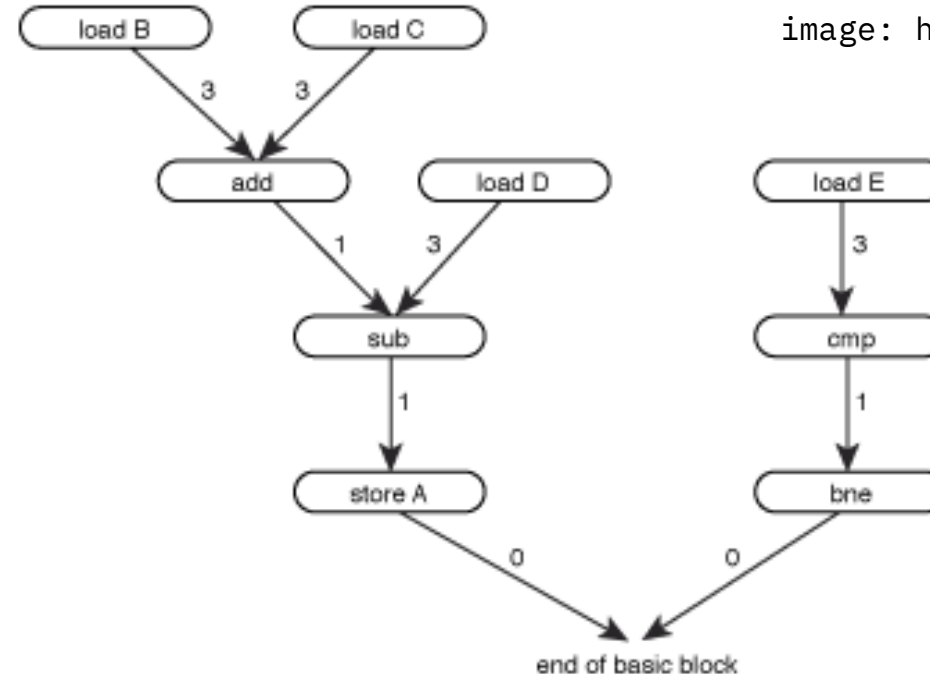


image: <http://bit.ly/2ERjDGY>

A Likely Schedule

```
load  b, r1  
load  c, r2  
load  d, r4  
load  e, r6  
add    r1, r2, r3  
sub    r3, r4, r5  
cmp    r6, 0  
store  r5, a  
bne    ...
```

Speculative Execution

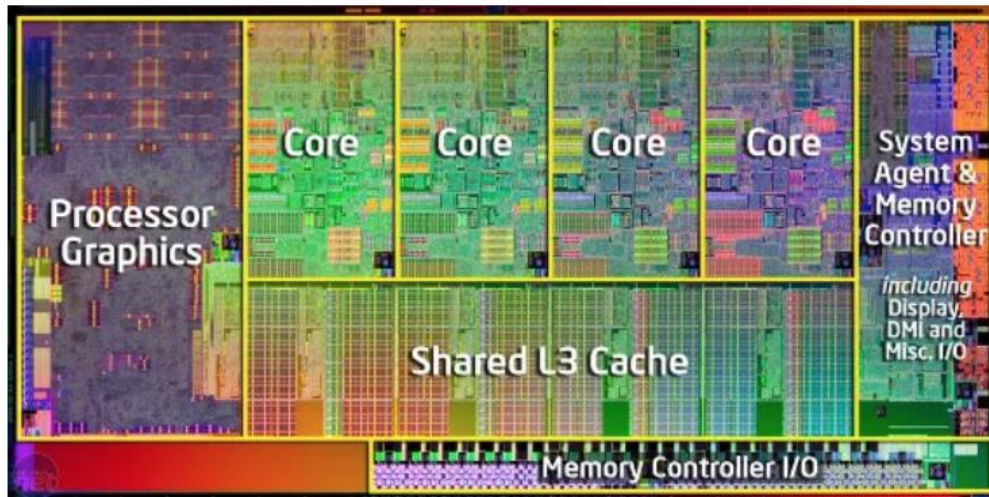
- This is a combination of:
 - branch prediction; plus
 - data-flow analysis
- Idea:
 - Perform some computations far ahead of time (i.e., speculatively)
 - This is done before we know they will be used...
 - ... although branch prediction helps increase the odds.
 - If results are not needed, then they are simply discarded.
- The observation here is that:
 - mainline code + speculative code can be evaluated in parallel

Spectre!

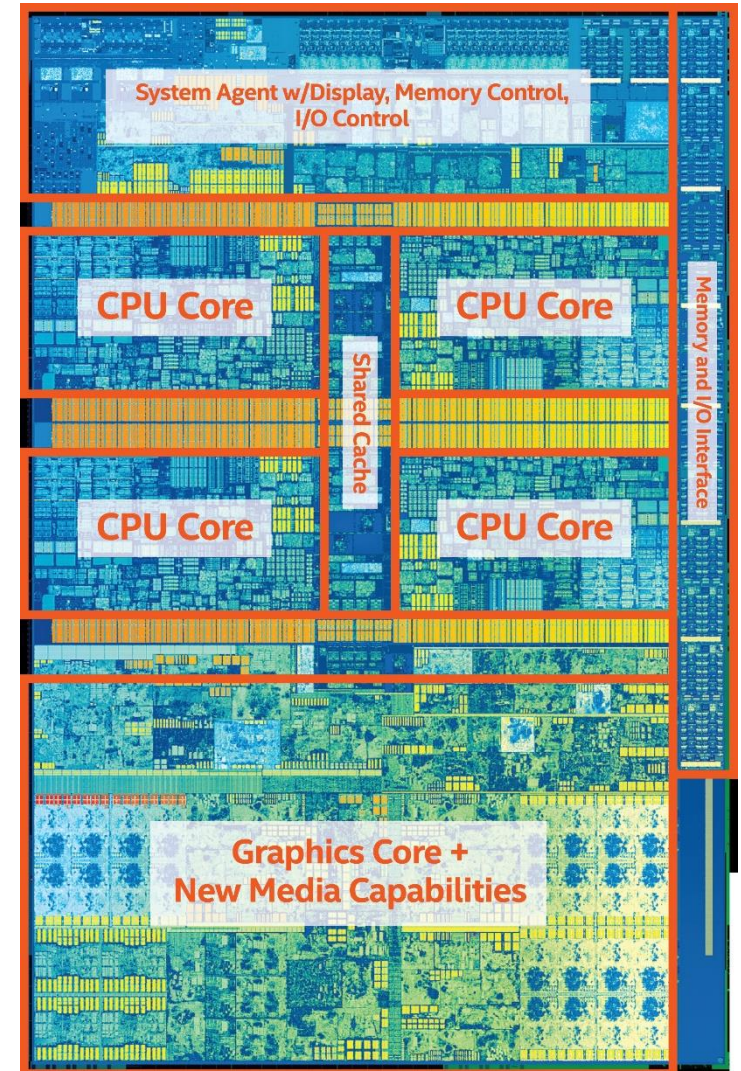
Meltdown!

Multicore

- Multiple cores == multicore
 - simply put more CPUs onto the same die!
- Different tasks can be executed on a different CPU cores (Multitasking)
- Real parallel execution
- One mechanism for this: **multithreaded code**




Images: <http://cnet.co/2EHdIRu>, <http://bit.ly/2BEzJ4V>





Multicore (Cont.)

- Each core is able to run an independent sequence of instructions from other cores
 - Some of the onchip DRAM cache is shared amongst cores (Level 2)
 - ... but some cache is not (Level 1)
 - More about this later in the course when we examine the memory hierarchy
- Intel also introduced **hyperthreading**
 - Term is a bit confusing as it is not under programmer control (like regular threads)
 - Processor itself is able to make a single core look like more than one
 - **This is done without the intervention from the software or from a software developer!**



Hyper-Threading Technology allows the CPU exposes two execution contexts per physical core.

Single- vs. Multi-threaded



Speedup Analysis

- Consider a simpler example: parallel (i.e., multiple cores)
- Question:
 - Given some program P ...
 - ... that takes time T to execute on a single core ...
 - ... what is the **greatest speedup possible** when using N processors/cores
- Assumption:
 - We have programs for which some fraction of the code can be into a parallel form
 - We name this fraction f such that $0 \leq f \leq 1$
 - Therefore, the fraction of the code which **cannot be parallelized** is $(1-f)$

Amdahl's Law

- This can be applied to evaluate any design or technical improvement


$$\text{speedup} = \frac{\text{execution time **before** enhancement}}{\text{execution time **after** enhancement}}$$

- Assumption here is that we can always characterize the fraction of code that can be enhanced (i.e., f)

Possible benefits from multiple cores

- If we have a single core:
 - Time to execute: T
- If the whole program is parallelizable (i.e., $f = 1.0$) and we have N cores:
 - New time to execute: $\frac{T}{N}$
- If part of the program is parallelizable ($0 < f < 1.0$), and we have N cores:
 - New time to execute: $T(1 - f) + \frac{Tf}{N}$

Possible benefits from multiple cores (Cont.)



Possible benefits from multiple parallel

- The ratio of **pre-improvement time** to **post-improvement time** is the **speedup** of the improvement
- $$\text{Speedup} = \frac{T}{T(1-f) + \frac{Tf}{N}} = \frac{1}{(1-f) + \frac{f}{N}}$$
- Note:
 - When f is small, speedup is negligible
 - As $N \rightarrow \infty$, speedup is bound by $1/(1-f)$
- These equations and observations are collectively known as **Amdahl's Law**

Possible benefits from multiple parallel (Cont.)




Image: Stallings (10th ed),
figure 2.4

Measures of computer performance

■ Warning:

- In the previous slides we used the symbol f to mean a fraction of code.
- We re-use the symbol f in what follows, but to mean a different kind of quantity
- Sorry, but that's the nature of the field

■ Recall:

- Processor's operations driven by the clock that runs at certain frequency f
- Time between clock pulses (cycle time): $\tau = \frac{1}{f}$
- Each processor instruction requires a fixed number of cycles
- Different types of instruction require different numbers of cycles

Given a specific program & architecture...

- I_C is the **instruction count**
 - number of machine instructions executed for that whole program until it runs to completion
- I_i is the count of instructions executed of type i
- CPI_{i_i} is the number of cycles needed to execute an instruction of type i

- Overall $CPI = \frac{\sum_{i=1}^n (CPI_i \times I_i)}{I_C}$

- Time to execute a given program:

$$T = I_C \times CPI \times \tau$$

Given a specific program & architecture...

- Suppose we had four instruction classes:

- ALU
- Memory (load, store)
- Branches (conditionals)
- Jumps (unconditional)

- Overall CPI =
$$\frac{(CPI_{ALU} \times I_{ALU}) + (CPI_{MEM} \times I_{MEM}) + (CPI_{BR} \times I_{BR}) + (CPI_J \times I_J)}{I_c}$$

Slight re-writing of CPI

- Our formulation of CPI does not make explicit the effects of memory latency
- We can re-write the equation to bring out this detail:
 - cycles to decode and execute instruction: p
 - number of memory references: m
 - ratio of memory-cycle time to processor-cycle time (i.e., how much slower or faster memory is compared to CPU): k
- $T = I_C \times [p + (m \times k)] \times \tau$

MIPS

- **Millions of instructions per second**

- Common measure
- (It is, however, a bit misleading)

- $$\text{MIPS rate} = \frac{I_c}{T \times 10^6} = \frac{f}{CPI \times 10^6}$$

- **Example:**

- 400 MHz processor
- Instruction type 1: Arithmetic & Logic (CPI = 1, 60% of instructions)
- Instruction type 2: Load/store from cache (CPI = 2, 18% of instructions)
- Instruction type 3: Branch (CPI = 4, 12% of instructions)
- Instruction type 4: Load/store from outside cache (CPI=8, 10% of instructions)
- $CPI = 0.6 + (2 \times 0.18) + (4 \times 0.12) + (8 \times 0.1) = 2.24$
- $\text{MIPS rate} = \frac{400 \times 10^6}{2.24 \times 10^6} \cong 178$



Any Questions?